

CSC 360

# Operating System Structures: **From Processes to Threads**


# Threads

---

- What is shareable with processes before threads?
- Motivation for threads

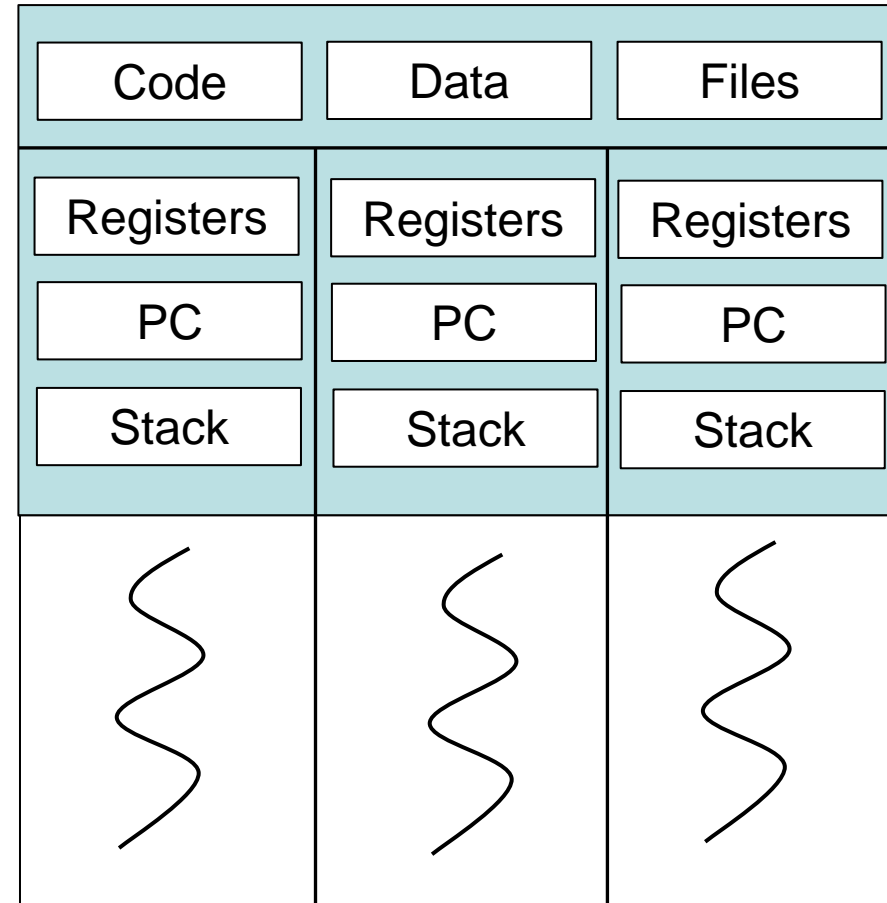
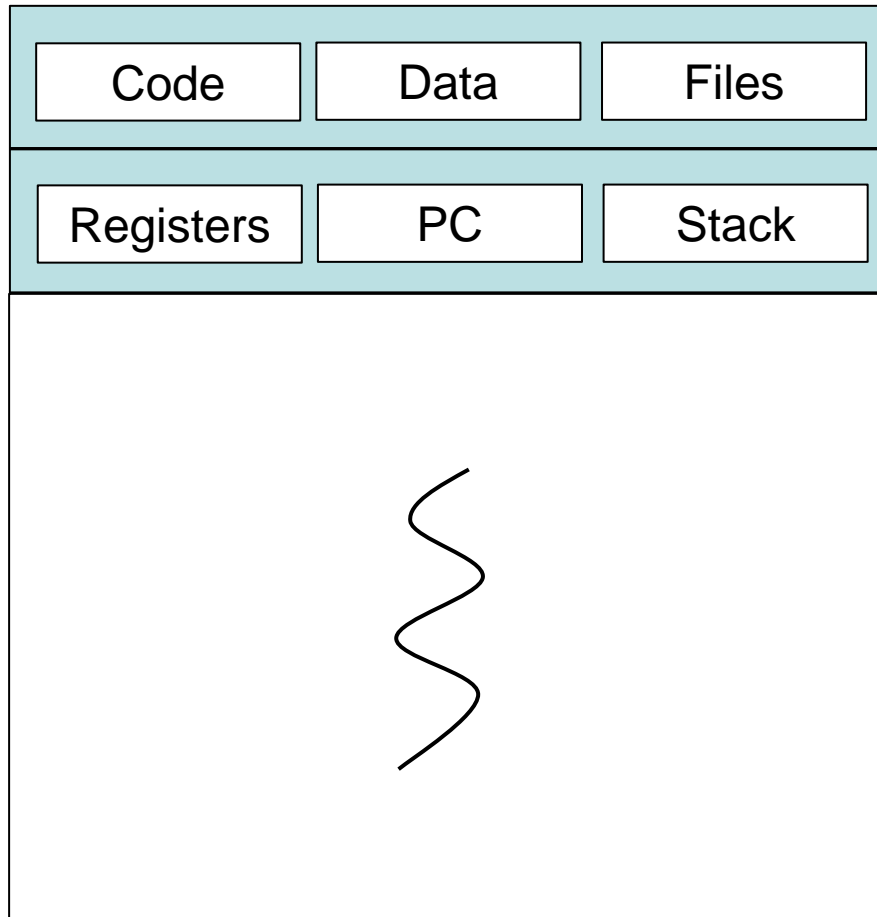
# Why Threads?

---



- **Thread**
  - short for **thread of control**
  - the sequence of executed instructions in a program, *representing a single path of execution*
- Many algorithms are **easier** to write (and maintain) with threads
- Some algorithms run **faster** in a threaded implementation.

# Single-thread vs. multithreaded process



# Process vs. Thread

---

- Processes have **separate address spaces**
  - Enforced by OS
  - *Sharing memory amongst processes requires OS intervention* (i.e., somewhat complex and expensive in time/CPU cycles)
- Threads exist **within** a process
  - A process may have **one or many threads**
  - Each of these threads **shares the same address space**
  - That is, *these threads by default share memory of their host process...*
  - ... **and this sharing does not require the involvement of the kernel!**
  - For threads this semester, we'll focus on POSIX threads (pthreads).


# Benefits of Using Threads

---

- Using threads is a natural way to achieve concurrency in programming
- Threads enable us to take advantage of multiprocessor systems,...
  - but they are just as useful on uniprocessor systems
- In many cases, a multithreaded solution is easier to **develop**, **understand**, and **debug** than its single-threaded counterpart
- By allowing multiple tasks to run concurrently, threads improve program responsiveness and resource utilization

# Motivation...

---

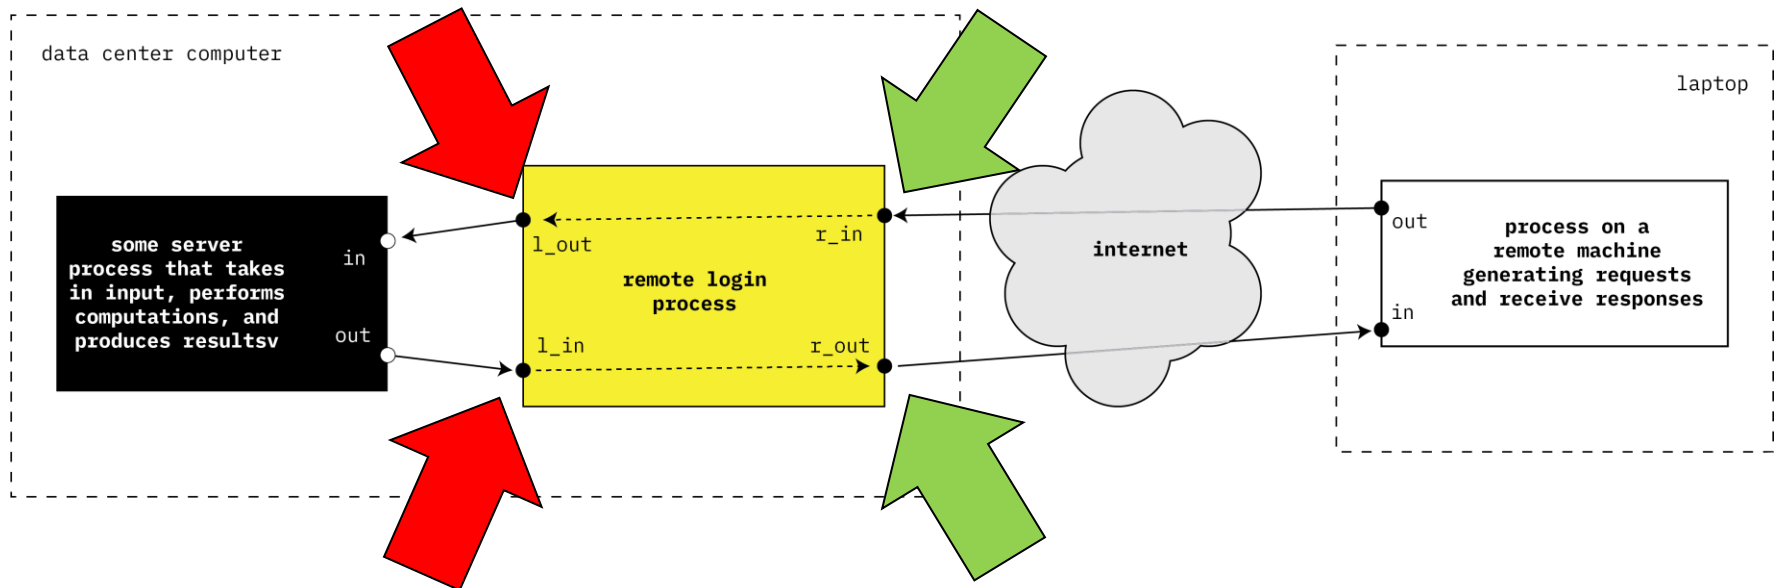


- **Server application**
  - Example: **remote login server**
  - Other examples: web server; database server

Describing action from the point of view of the **remote login process**.  
The names of the ports are from the point of view of this process.

(B) Requests intended for server processing **are sent out here.**

(A) Incoming messages from remote machine for server processing **arrive here**



(C) Results produced by server processing **are sent out here.**

(D) Outgoing messages containing the results from server processing **are sent out here.**



# Purpose of a remote login server

---

- This server (or *daemon*) acts as the go-between:
  - the computer outside the server machine, and ...
  - the processes within the server machine
- Remote-login server **receives** incoming messages from remote computer, and **sends** them to a process within the server computer
- Remote-login server **receives** responses from process within the server computer, and **sends** them to the remote computer
- Also:
  - The **r\_in** and **l\_out** ports only have messages when the remote computer or server machine have something to send (i.e., remote login server may need to wait for message on those incoming message ports)
  - The **l\_in** and **r\_out** ports have outgoing message queues that might be full (i.e., remote login server may need to wait for outgoing message ports to have room for new messages)

# One implementation approach (bad)

---

- Remote login server (daemon) is organized around a **polling loop**
  - For each loop iteration, code examine each of the four ports
  - If a port requires servicing, then something is done
  - Otherwise return to top of loop
- **Observation:** Given the speed of CPUs relative to the frequency of server events, most loop iterations will do no processing (i.e., nothing to do) ....
  - Therefore, CPU cycles are wasted
- Polling is often used as a first approach to solving such coding problems...
  - ... but the wasted CPU cycles can become a problem as more and more such daemons exists in a computer

# Better implementation approach

---

- We still use a loop, but its contents are very different (***but still without threads...***)
- Code in loop uses system calls such that the OS:
  - **Suspends** the remote-login-daemon's process when there is no work (i.e., no messages to relay from port to port)
  - **Reawakens** the remote-login-daemon's process when there is a message to relay to a port
- **This is far, far better than a straight polling loop**
  - When the daemon has “nothing to do”, the process does not consume CPU cycles
  - When the daemon does have something to do, it is normally
    - (1) moving messages from **r\_in** to **l\_out**, or
    - (2) moving messages from **l\_in** to **r\_out**
- One issue, however, with respect to the design of our code:
  - If our process possesses a single thread, then the code of the loop **must be carefully written to interleave actions (1) and (2)**

# Life Without Threads

Is the purpose of the code immediately apparent?

```
logind(int r_in, int r_out, int l_in, int l_out) {
    fd_set in = 0, out;
    int want_l_write = 0, want_r_write = 0;
    int want_l_read = 1, want_r_read = 1;
    int eof = 0, tsize, fsize, wret;
    char fbuf[BSIZE], tbuf[BSIZE];

    fcntl(r_in, F_SETFL, O_NONBLOCK);
    fcntl(r_out, F_SETFL, O_NONBLOCK);
    fcntl(l_in, F_SETFL, O_NONBLOCK);
    fcntl(l_out, F_SETFL, O_NONBLOCK);

    while(!eof) {
        FD_ZERO(&in);
        FD_ZERO(&out);
        if (want_l_read) FD_SET(l_in, &in);
        if (want_r_read) FD_SET(r_in, &in);
        if (want_l_write) FD_SET(l_out, &out);
        if (want_r_write) FD_SET(r_out, &out);

        select(MAXFD, &in, &out, 0, 0);

        if (FD_ISSET(l_in, &in)) {
            if ((tsize = read(l_in, tbuf, BSIZE)) > 0)
            {
                want_l_read = 0;
                want_r_write = 1;
            } else {
```

# Life Without Threads

Is the purpose of the each “if” statement immediately apparent?

```
    eof = 1;
}
if (FD_ISSET(r_in, &in)) {
    if ((fsize = read(r_in, fbuf, BSIZE)) > 0)
    {
        want_r_read = 0;
        want_l_write = 1;
    } else
        eof = 1;
}

if (FD_ISSET(l_out, &out)) {
    if ((wret = write(l_out, fbuf, fsize)) == fsize)
    {
        want_r_read = 1;
        want_l_write = 0;
    } else if (wret >= 0)
        tsize -= wret;
    else
        eof = 1;
}

if (FD_ISSET(r_out, &out)) {
    if ((wret = write(r_out, tbuf, tsize)) == tsize)
    {
        want_l_read = 1;
        want_r_write = 0;
    } else if (wret >= 0)
        tsize -= wret;
    else
        eof = 1;
}
}
```

```

logind(int r_in, int r_out, int l_in, int l_out) {
    fd_set in = 0, out;
    int want_l_write = 0, want_r_write = 0;
    int want_l_read = 1, want_r_read = 1;
    int eof = 0, tsize, fsize, wret;
    char fbuf[BSIZE], tbuf[BSIZE];

    fcntl(r_in, F_SETFL, O_NONBLOCK);
    fcntl(r_out, F_SETFL, O_NONBLOCK);
    fcntl(l_in, F_SETFL, O_NONBLOCK);
    fcntl(l_out, F_SETFL, O_NONBLOCK);

    while(!eof) {
        FD_ZERO(&in);
        FD_ZERO(&out);
        if (want_l_read) FD_SET(l_in, &in);
        if (want_r_read) FD_SET(r_in, &in);
        if (want_l_write) FD_SET(l_out, &out);
        if (want_r_write) FD_SET(r_out, &out);

        select(MAXFD, &in, &out, 0, 0);

        if (FD_ISSET(l_in, &in)) {
            if ((tsize = read(l_in, tbuf, BSIZE)) > 0)
            {
                want_l_read = 0;
                want_r_write = 1;
            } else {

```

**l\_in:** Local input      **l\_out:** Local output  
**r\_in:** Remote input    **r\_out:** Remote output

Initially, we want to **read** from the local terminal & remote machine;  
 We do not write anything until we have data

Set Non-Blocking Mode; This ensures that **read()** and **write()** do not block execution

Monitoring I/O using **select()** system call;  
**continuously monitors the file descriptors** until at least one file descriptor is ready for reading or writing.

### Read Data from Local Terminal

If **l\_in** has data, we read it into **tbuf**;

1. We pause further reading (**want\_l\_read = 0**)
2. and prepare to write this data to the remote machine (**want\_r\_write = 1**)

```

    eof = 1;
}
if (FD_ISSET(r_in, &in)) {
    if ((fsize = read(r_in, fbuf, BSIZE)) > 0)
    {
        want_r_read = 0;
        want_l_write = 1;
    } else
        eof = 1;
}

if (FD_ISSET(l_out, &out)) {
    if ((wret = write(l_out, fbuf, fsize)) == fsize)
    {
        want_r_read = 1;
        want_l_write = 0;
    } else if (wret >= 0)
        tsize -= wret;
    else
        eof = 1;
}

if (FD_ISSET(r_out, &out)) {
    if ((wret = write(r_out, tbuf, tsize)) == tsize)
    {
        want_l_read = 1;
        want_r_write = 0;
    } else if (wret >= 0)
        tsize -= wret;
    else
        eof = 1;
}
}
}

```

## Read Data from Remote Machine

If **r\_in** has data, we read it into **fbuf**

1. We pause further reading from the remote (**want\_r\_read = 0**) and
2. prepare to write this data to the local terminal (**want\_l\_write = 1**).

## Write Data to Local Terminal

If **l\_out** is ready for writing, **we send the data received from the remote machine.**

Once the full message is sent, **we resume reading from the remote.**

## Write Data to Remote Machine

If **r\_out** is ready for writing, we send the data received from the local terminal.

Once the full message is sent, we resume reading from the local terminal.

# Better implementation approach

---

- **Multithreading:**
  - Our process will now consist of two separate threads
  - Both threads will share the processes global (i.e., program scope) variables/data...
  - ... but each thread will have its own local variables.
- **The important idea:**
  - One thread handles the message transfer from **r\_in** to **l\_out**...
  - ... while the other thread handles the message transfer from **l\_in** to **r\_out**
- (We will see later how to start these threads, but for now let us look at one possible version of their code)



# Life With Threads

---

```
incoming(int r_in, int l_out) {
    int eof = 0;
    char buf[BSIZE];
    int size;

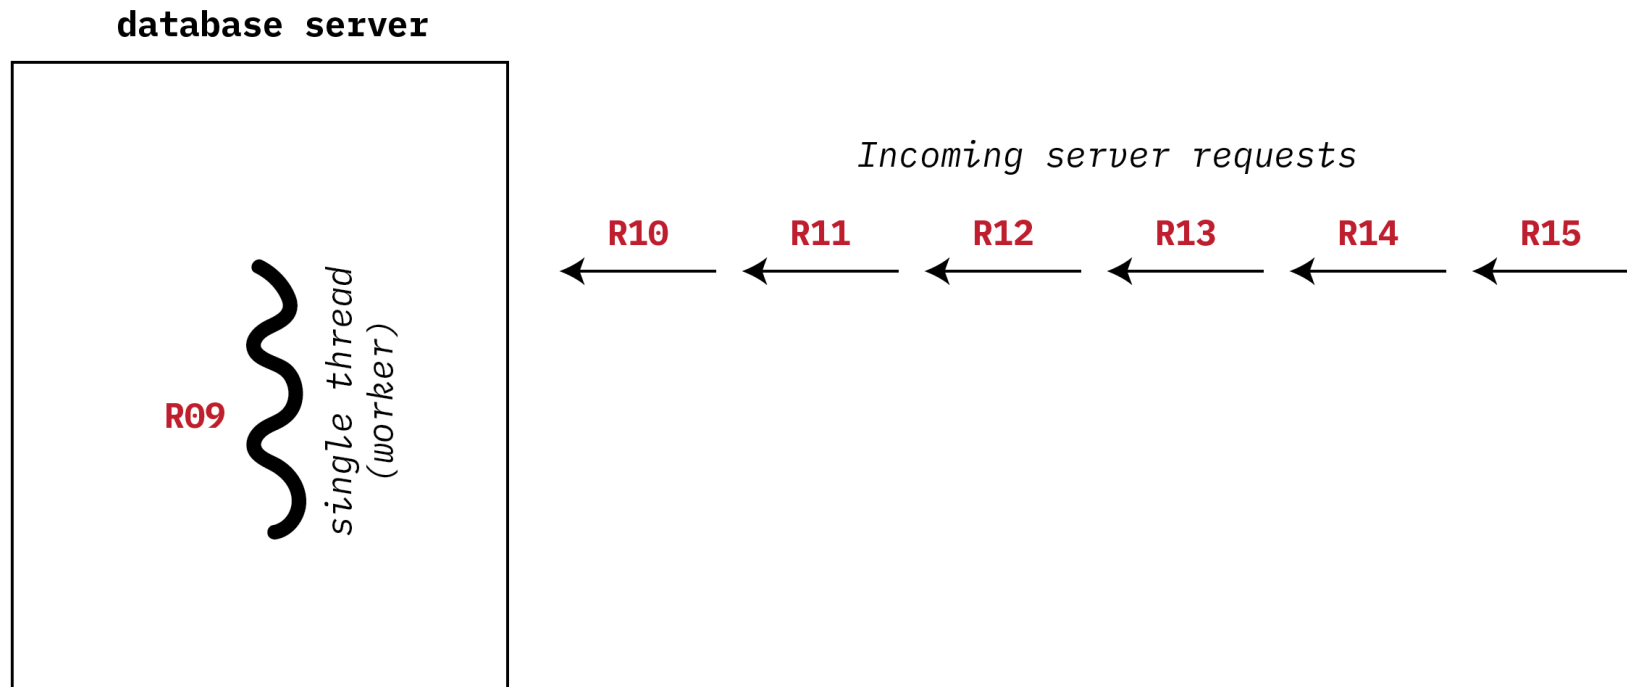
    while (!eof) {
        size = read(r_in, buf, BSIZE);
        if (size <= 0)
            eof = 1;
        if (write(l_out, buf, size) <= 0)
            eof = 1;
    }
}
```

```
outgoing(int l_in, int r_out) {
    int eof = 0;
    char buf[BSIZE];
    int size;

    while (!eof) {
        size = read(l_in, buf, BSIZE);
        if (size <= 0)
            eof = 1;
        if (write(r_out, buf, size) <= 0)
            eof = 1;
    }
}
```

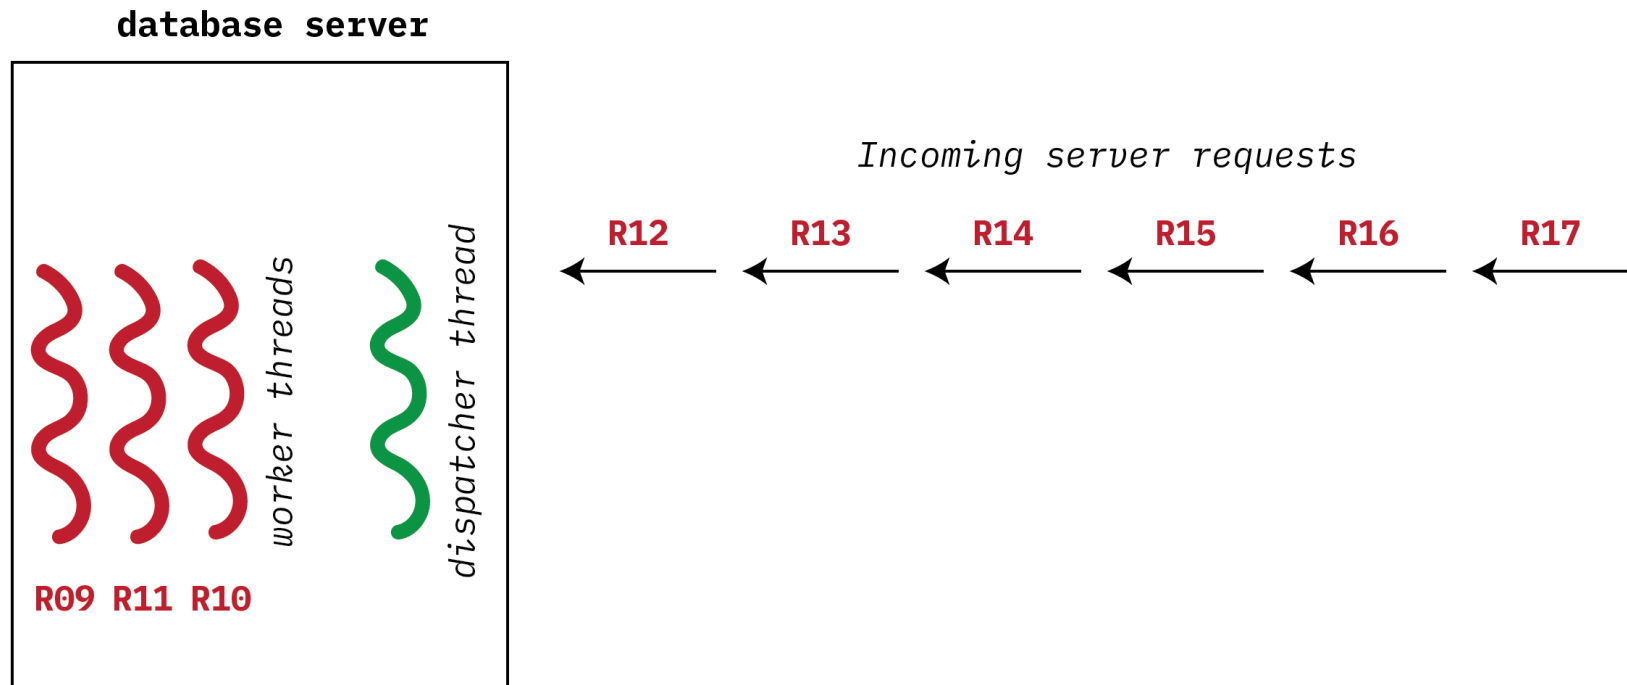
# Single-Threaded Database Server

---



# Multithreaded Database Server

---



# Benefits of threads

---

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process (open files, program-scope variables) which is easier to use than shared memory or message passing between processes
- **Economy** – thread creation is “cheaper” than process creation; thread switching has lower overhead than context switching
- **Scalability** – program in the process can be written to take advantage of multiprocessor architectures by adding and removing threads and processors become more or less available

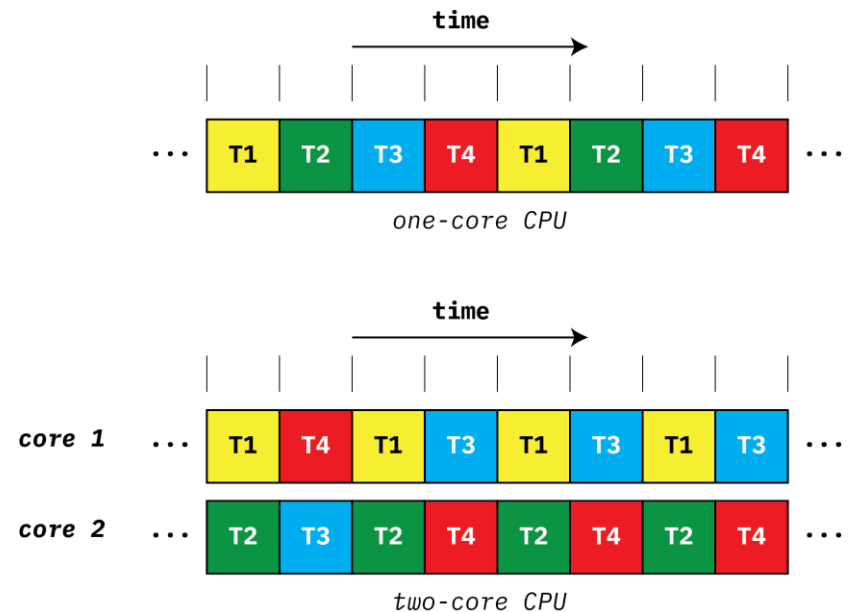
# Multicore Programming

---

- Multicore or multiprocessor systems do, however, put pressure on programmers
- Programming challenges include:
  - **Dividing up activities**
  - **Balancing activities**
  - **Data splitting**
  - **Data dependency**
  - **Testing and debugging**
- Parallelism implies a system can perform more than one task simultaneously
- Concurrency supports more than one task making progress
  - Single processor / core, scheduler providing concurrency

# Multicore Programming (Cont.)

- Types of parallelism
  - **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
  - **Task parallelism** – distributing threads across cores, each thread performing unique operation
- Hardware architecture support for software threads continues to increase over time
  - CPUs have **cores** as well as **hardware threads**
  - Example: Intel Xeon Platinum 8592+ processors has 64 cores (price: US\$12K)
  - **Note: hardware threads are *not the same* as the software threads we're examining in this course.**



# pthread (POSIX threads)

---

- The creation and management of threads is normally done through a **thread library**
- **We will use POSIX threads** (also known as pthreads)
  - Which means we will be using the pthread library installed with on `jhub-cos`
- In what follows we will look at some examples of pthread library calls
  - thread **creation**
  - **parameterizing** threads
  - thread **attributes**
  - and more...
- More specific details on pthreads will eventually be provided as part of assignments and tutorials

# Creating a POSIX Thread

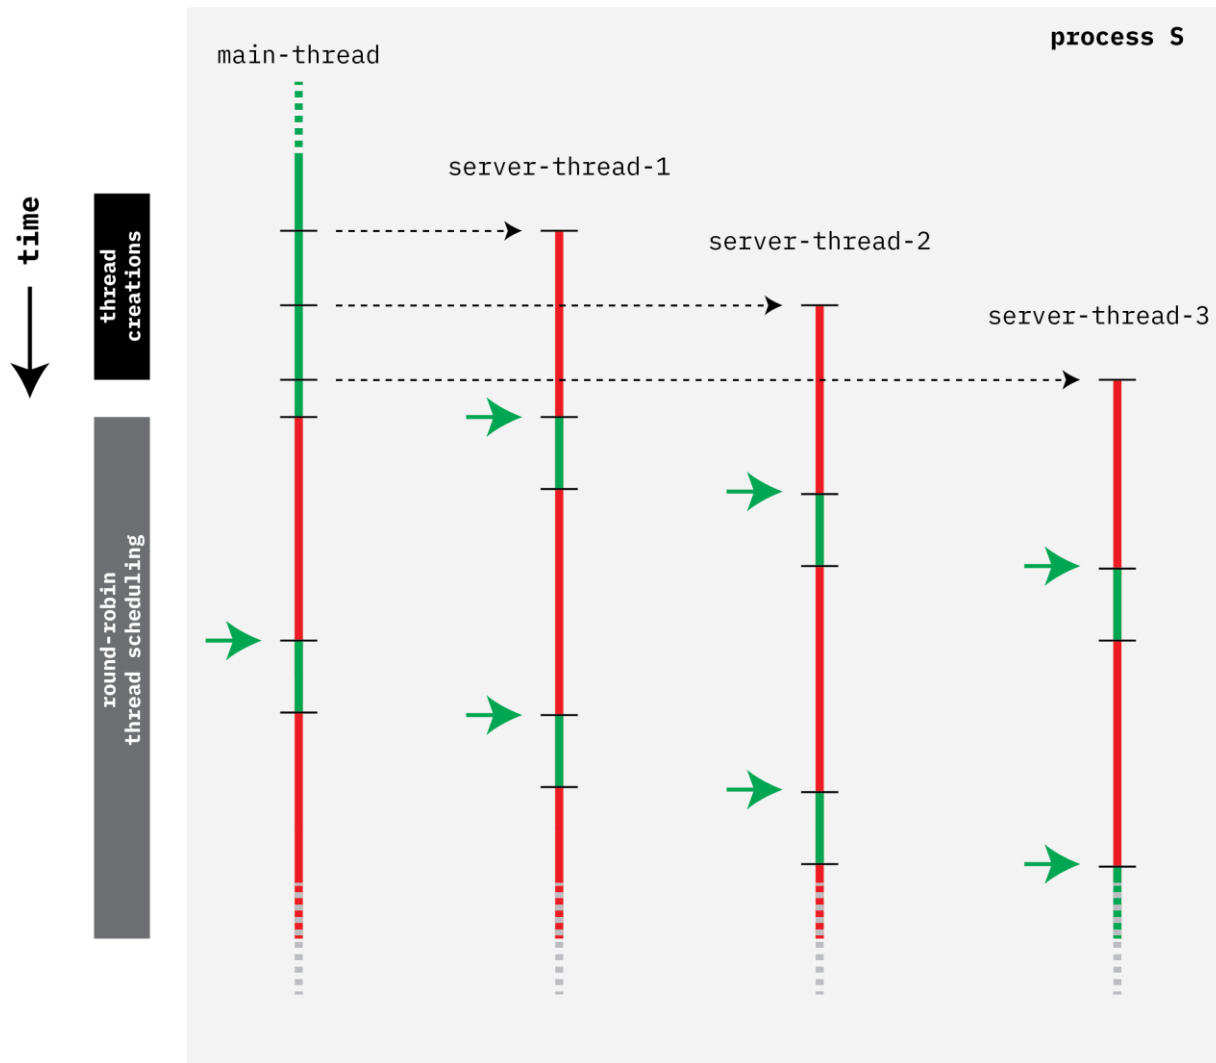
---

```
void start_servers( ) {
    pthread_t thread;
    int i;
    for (i=0; i<nr_of_server_threads; i++)
        /* Making a biiiiiig assumption below, that thread
         * creation always succeeds. More robust code
         * must check value returned from pthread_create.
         */
        pthread_create(&thread, /* pointer to thread ID */
                      0, /* default attributes */
                      server, /* start routine */
                      argument); /* pointer to any argument */
}

void *server(void *arg) { /* the worker thread */
    for (;;) {
        /* get and handle request */
    }
}
```



# Threads in one process



# (a wee complication)

---

```
rlogind(int r_in, int r_out, int l_in, int l_out)
{
    pthread_t in_thread, out_thread;

    pthread_create(&in_thread,
        0,
        incoming,
        r_in, l_out);          /* Sadly, cannot do this... */

    pthread_create(&out_thread,
        0,
        outgoing,
        l_in, r_out);          /* Can't do this, either... */

    /* How do we wait till both threads are done? */
}
```

# Multiple Arguments

---

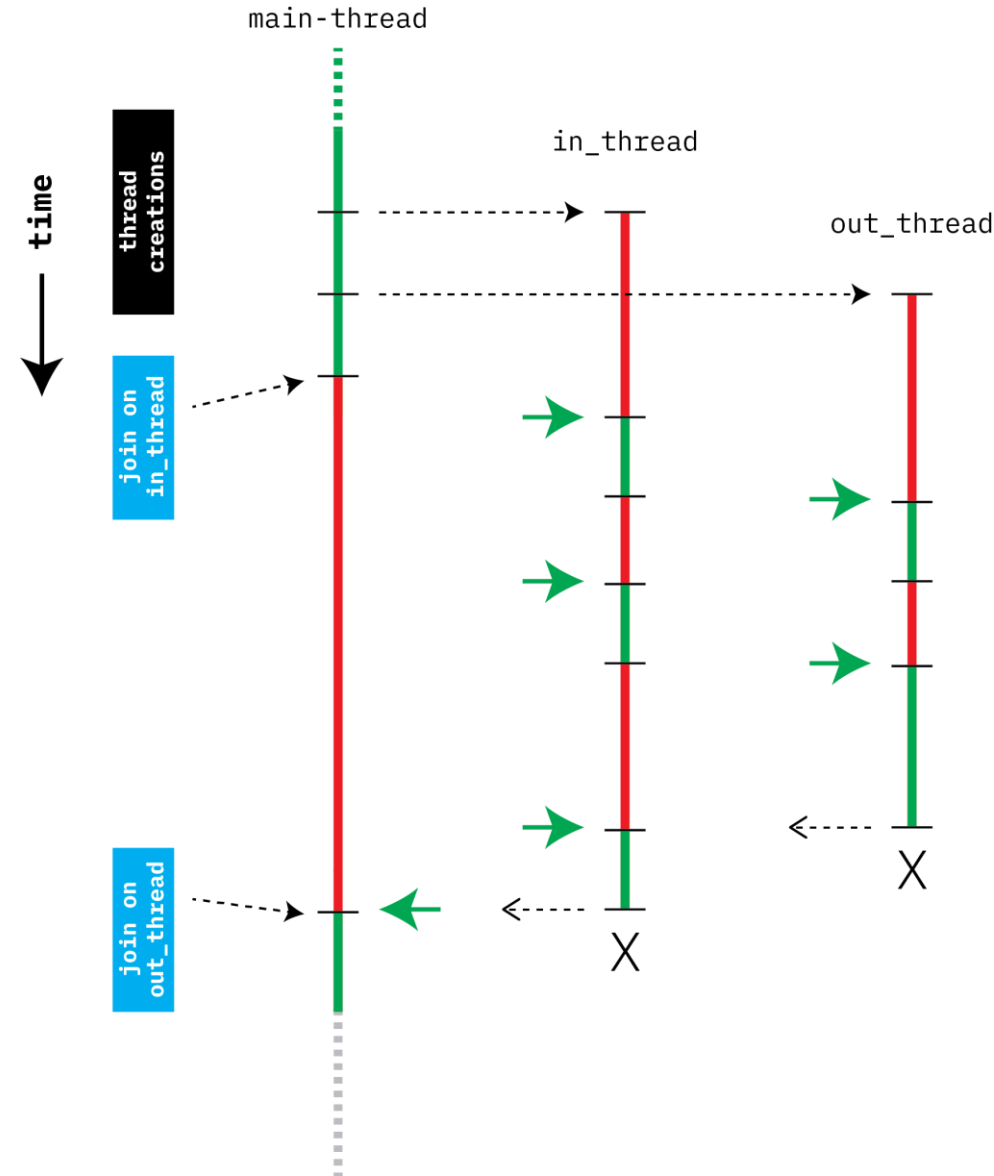
```
typedef struct {  
    int first, second;  
} two_ints_t;  
  
rlogind(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
  
    two_ints_t in = {r_in, l_out}, out = {l_in, r_out};  
  
    pthread_create(&in_thread, 0, incoming, &in);  
    pthread_create(&out_thread, 0, outgoing, &out);  
}
```

# When are the threads done?

---

```
void rlogind(int r_in, int r_out, int l_in, int l_out) {  
    pthread_t in_thread, out_thread;  
    two_ints_t in={r_in, l_out}, out={l_in, r_out};  
  
    pthread_create(&in_thread, 0, incoming, &in);  
    pthread_create(&out_thread, 0, outgoing, &out);  
  
    pthread_join(in_thread, 0);  
    pthread_join(out_thread, 0);  
}
```

- `pthread_join()` blocks execution until both threads complete their work.
- Ensures that both threads finish before `rlogind` (mainline code) exits.



```
void rlogind(int r_in, ...
    ...

    pthread_create(&in_thread, ...
    pthread_create(&out_thread, ...

    pthread_join(in_thread, 0);
    pthread_join(out_thread, 0);
}
```

Notice that even though **out\_thread** completes before **in\_thread**, the **main-thread** is blocked until it joins with **in\_thread**.

# Termination

---

- Careful use of several calls (signatures shown below)
  - #1 is completely different from regular process `exit(1)`
  - #2 may be needed to keep the compiler happy
  - #3 really only needed if the thread creator must explicitly wait for the created thread to complete its work.

```
/* 1 */ pthread_exit((void *) value);  
  
/* 2 */ return((void *) value);  
  
/* 3 */ pthread_join(thread, (void **)&value);
```

# Detached Threads

---

```
void start_servers( ) {  
    pthread_t thread;  
    int i;  
    for (i=0; i<nr_of_server_threads; i++) {  
        pthread_create(&thread, 0, server, 0);  
        pthread_detach(thread);  
    }  
    ...  
}  
  
server( ) {  
    ...  
}
```

# Thread Attributes

---

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
  
...  
/* establish some attributes */  
...  
  
pthread_create(&thread, &thr_attr, startroutine, arg);  
  
...  
pthread_attr_destroy(&thr_attr);
```



# Stack Size



---

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
pthread_attr_setstacksize(&thr_attr, 20*1024*1024);  
  
...  
  
pthread_create(&thread, &thr_attr, startroutine, arg);  
  
...  
pthread_attr_destroy(&thr_attr);
```

**View of the process address space showing multiple local stacks.**

Note: For simplicity, kernel portion of the address space is not shown.

all threads possess  
the same code!



# User Threads and Kernel Threads

---

- **User threads**
  - Management performed by code in a user-level threads library
  - Kernel **is not aware of threads** (that is, any system call from any thread appears to the thread **as if it comes from the whole process**)
- Three primary thread libraries:
  - Traditional POSIX Pthreads
  - Windows threads
  - Java threads (but this requires a Java VM)
- **Kernel threads**
  - Supported by code in the kernel
  - Kernel **is aware of system calls** (including blocking ones) made by thread
- Examples: virtually all general purpose operating systems, including
  - Windows
  - Solaris
  - **Linux (i.e., pthread)**
  - Tru64 UNIX
  - macOS


# Multithreading Models

---

- Many-to-One
  - Many user threads to one kernel thread
- One-to-One
  - One user thread to one kernel thread
- Many-to-Many
  - M user threads to N kernel threads

# Many-to-One


- Many user-level threads mapped to single kernel thread
- **One thread blocking causes all to block**
- Multiple threads may not run in parallel on a multicore system because only one may be in kernel at a time
- Although this was the very first widely-implemented form of threading, **few systems currently use this model**
- Examples:
  - Solaris Green Threads
  - GNU Portable Threads



# One-to-One


---

- Each user-level thread maps to kernel thread
- **Creating a user-level thread creates a kernel thread**
- More physical concurrency possible than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
  - Windows NT
  - Linux
  - Solaris 9 and later




# Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
  - (Also known as an M to N -- or **M:N model**)
- Allows the operating system to **create a sufficient number of kernel threads**
- Solaris prior to version 9
- Windows with the ThreadFiber package



# Two-level Model


- Similar to M:N, except that it allows a single user thread within a process to be bound to single kernel thread
  - That is, all threads are in the same process...
  - ... but some are M:N, some are 1:1
- Examples
  - IRIX
  - HP-UX
  - Tru64 UNIX
  - Solaris 8 and earlier





# Scheduler Activations

- Both M:N and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
  - Appears to be a virtual processor on which process can schedule user thread to run
  - Each LWP attached to kernel thread
  - How many LWPs to create?
- Scheduler activations provide **upcalls**
  - a communication mechanism from the kernel to the upcall handler in the thread library
  - This communication allows an application to maintain the correct number kernel threads



# Linux Threads

---

- Linux refers to them as **tasks** rather than **threads**
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
  - Flags control behavior
- `struct task_struct` points to process data structures (shared or unique)

flag	meaning
<code>CLONE_FS</code>	File-system information is shared.
<code>CLONE_VM</code>	The same memory space is shared.
<code>CLONE_SIGHAND</code>	Signal handlers are shared.
<code>CLONE_FILES</code>	The set of open files is shared.

# Summary

---

- Threads are an important structural mechanism for programmatic expression of concurrency
  - Easier to write an application using multiple threads (i.e. compared with using multiple processes)
  - Threads can be implemented on top of a systems process model
- There are many different realizations of threads
  - Our course will focus on POSIX threads (pthreads)
  - Different models represented different relationships between user-level threads and kernel-level threads
- Creating threads is (oddly enough) the easy part!
  - **Ensuring threads collaborate together in safe ways is the hard part.**
  - We look next at this hard part...